A DB file in file on hard disk

> this experiment yields what happens when the number I/O with the db file is minimized (was done by others of course)
> we use this to show how our approach behaves in this context. Hypothesis: overhead we see is due to unsmart LRU not due to „db file lives outside the RAM“). Hopefully, we do far more better (i.e., „3“ has far better latency/ throughput) since smarter hot/cold data fetch/eviction

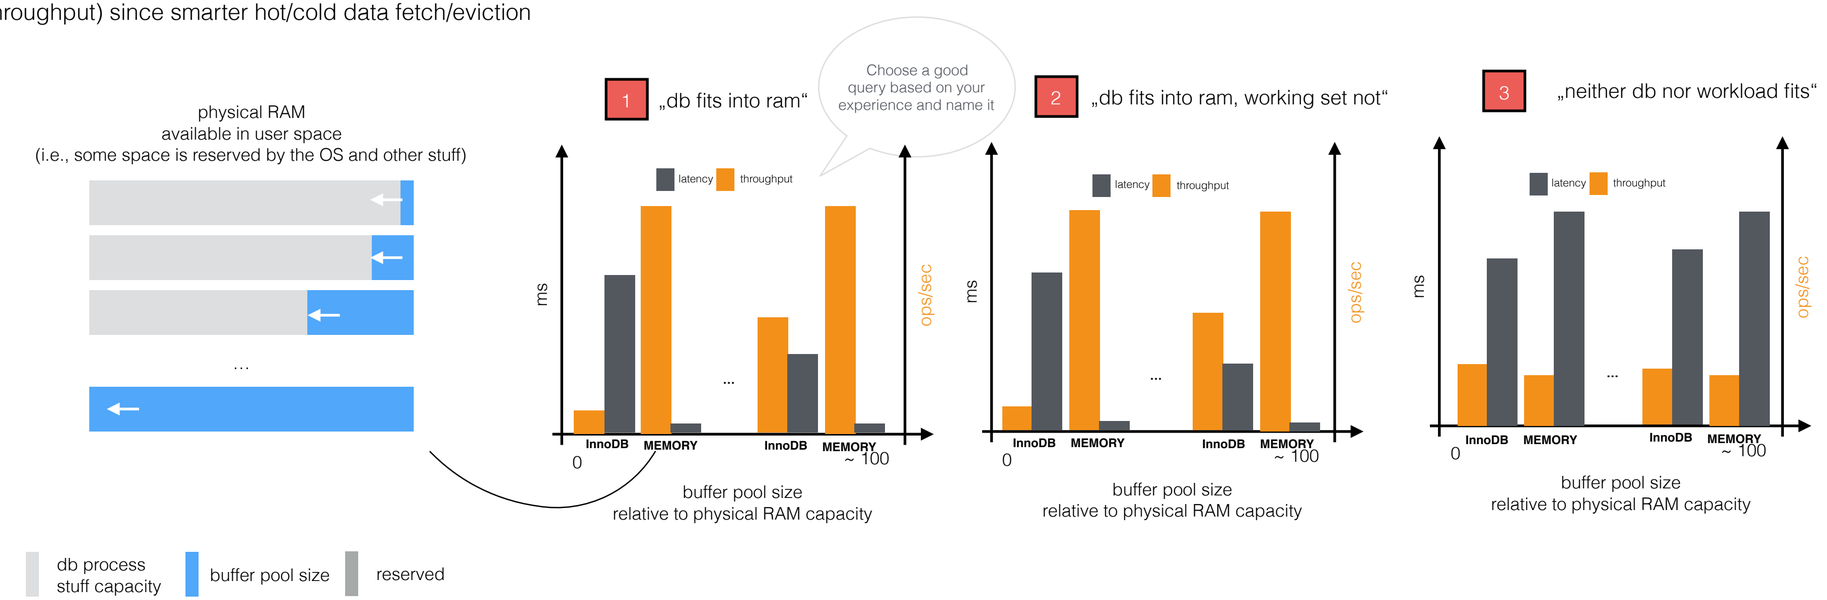

Allows or denies statement

The feature „tuple id“ abstraction is the issue.

B DB file in file on ram-disk (which is mapped into main memory!), DB *must* actually fit into main memory here

> this experiment yields the plain „overhead“ for buffer manager indirection (disk I/O is removed)
> we will use this to compare our allocator against this. Hypothesis: Reason is „tuple pair“ abstraction we have less indirection due to moving „tuple“ abstraction into RAM. Hence, better throughput and lower latency

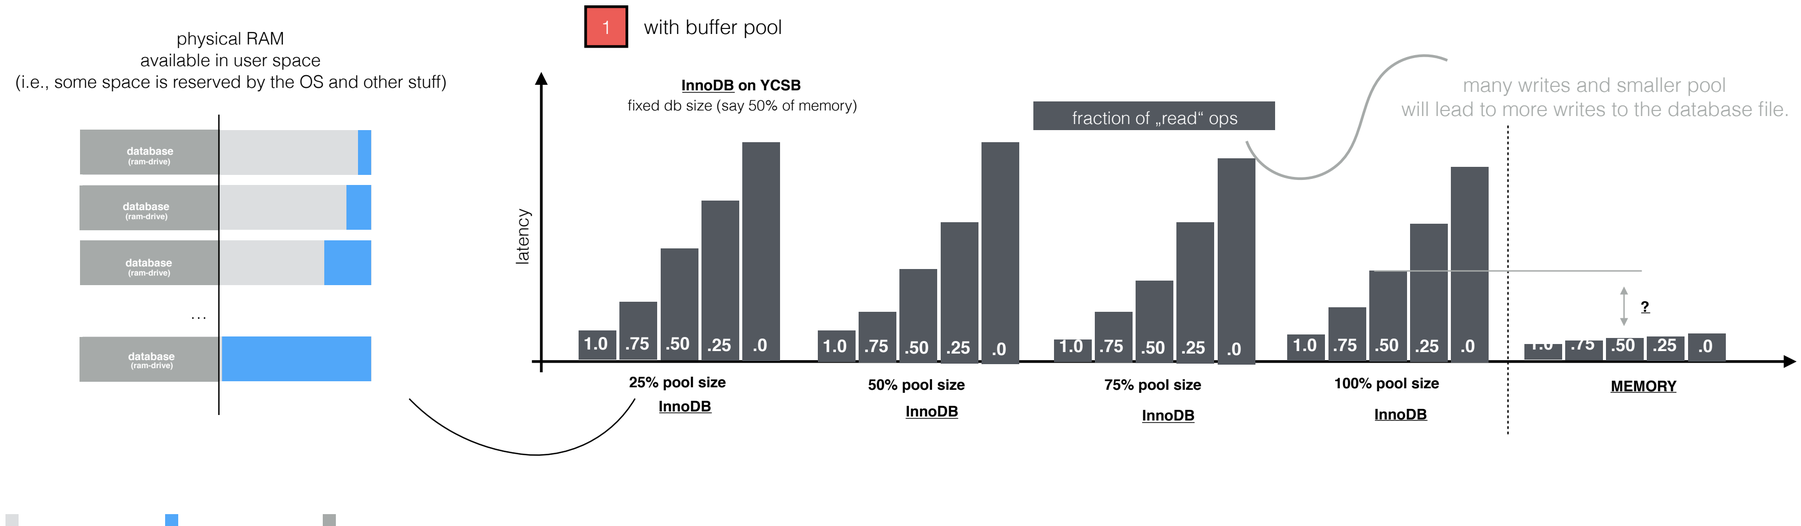

Allows or denies statement

The feature „persistent db is not in RAM“ is not the issue.